

AI at Work: Three Industries, One Month

Issue #1 · August 2026 · A monthly look at commercial AI in action

If you blinked this month, you missed a lot. AI didn't just get smarter in July and August, it got *deployed*. Enterprises pushed AI agents past the pilot stage, hospitals kept clearing new AI diagnostic tools at a record pace, and the tools developers use to build software kept reshaping how fast products actually ship.

This issue looks at three industries where AI isn't a future promise, it's already changing how decisions get made and work gets done. No hype, just what's actually happening and what it means for the businesses living through it.

IN THIS ISSUE

- **Enterprise Operations** — AI agents move from pilot to production, and the numbers on what's actually working
- **Healthcare** — A milestone FDA clearance and the steady march of AI into clinical decision-making
- **Software Development** — The coding assistant market matures, and terminal-native agents pull ahead

INDUSTRY 01

Enterprise Operations: Agents Grow Up

For the last two years, "AI agent" mostly meant an impressive demo. That's changing fast. This month, **OpenAI launched Presence**, a dedicated enterprise platform for deploying AI agents that can answer questions, use internal company systems, take approved actions, and know when to hand off to a human. It's a telling shift: the company that popularized the chatbot is now selling infrastructure for agents that actually do work, not just talk about it.


OpenAI Presence

Google Gemini Enterprise

Cognizant EMEA AI Unit

Oracle OCI Enterprise AI

Consulting giant **Cognizant** launched a dedicated EMEA AI unit in late July aimed squarely at a hard truth: most enterprise AI agent pilots never make it to production. Meanwhile, **Oracle** added day-zero support for NVIDIA's newest agent-optimized model to its cloud platform, betting that businesses want more flexibility in which "brain" powers their agents, not just faster versions of the same one.



Source: Gartner 2026 CIO and Technology Executive Survey

ANALYSIS: WHY THIS MATTERS

Only 17% of organizations have fully deployed AI agents today, but over 60% expect to within two years. That gap between ambition and reality is exactly why products like Presence and consulting practices like Cognizant's new unit exist: the technology has outpaced the organizational muscle needed to deploy it safely and reliably. Gartner also predicts over 40% of agentic AI projects will be cancelled by 2027, a reminder that the businesses winning here aren't the ones moving fastest, but the ones building proper

governance and monitoring alongside the AI itself. For decision-makers, the takeaway is blunt: an AI agent strategy without a deployment and oversight strategy is just an expensive pilot.

INDUSTRY 02

Healthcare: A New Kind of FDA Clearance

In late June, clinical AI company **UpDoc** received what's being described as the first FDA clearance for a Software as a Medical Device with a patient-facing large language model, meaning an AI system cleared to interact directly with patients using conversational AI, not just assist a clinician behind the scenes. It's a meaningful line to cross: most AI medical devices to date have been diagnostic aids reviewed by a doctor, not tools patients talk to directly.

UpDoc Inc.

Doctronic


ChatGPT Health

Viz.ai

1,300+

AI-enabled medical devices now cleared by the FDA, up from 950 just two years ago

This builds on a broader trend: Utah's pilot with Doctronic for autonomous AI prescription refills and OpenAI's own ChatGPT Health launch earlier this year both point the same direction, AI moving from the radiology back office into direct, patient-facing roles. Radiology still dominates FDA clearances, but cardiovascular and neurology applications are catching up.



Source: FDA AI-Enabled Medical Device List, tracked via Bipartisan Policy Center and IntuitionLabs

ANALYSIS: WHY THIS MATTERS

Healthcare has historically been one of the most conservative industries for AI adoption, and for good reason: mistakes have real human cost. The steady rise in FDA clearances shows regulators finding a workable rhythm with AI vendors rather than either rubber-stamping or blocking innovation. But patient-facing clearances like UpDoc's raise the stakes: the industry is no longer just asking "does this AI give a good answer," but "can this AI be trusted to interact with a vulnerable person directly." For healthcare businesses, the opportunity is real, faster diagnostics, better access in underserved areas, but so is the responsibility to get monitoring and human oversight right from day one.

INDUSTRY 03


Software Development: The Terminal Strikes Back

The AI coding assistant market has quietly become one of the fastest-moving corners of enterprise software. This year, **Claude Code**, Anthropic's terminal-native coding agent, pulled ahead on independent benchmarks, hitting roughly 80.8% on SWE-bench Verified and 89.1%

on the newer Terminal-Bench 2.1 suite as of this month, the strongest results of any tool tested. Unlike editor plugins, it works directly in a developer's command line: reading a codebase, editing files, running tests, and even committing code with minimal hand-holding.

- Claude Code (Anthropic)
- GitHub Copilot
- Cursor
- Amazon Q Developer

Competition is intensifying, not slowing down. GitHub Copilot moved to usage-based "AI Credit" billing in June, and its Vision capabilities reached general availability this month. Cursor, now backed by SpaceX after an acquisition earlier this year, keeps pushing multi-file, IDE-native editing. Meanwhile, pricing across the board keeps shifting as providers race to make agentic coding affordable enough for daily, all-day use rather than occasional help.



Source: Nipralo Technologies market analysis, 2026

ANALYSIS: WHY THIS MATTERS

The AI coding tools market has more than doubled in under two years, and roughly 90% of professional developers now use one of these tools daily. This isn't a novelty adoption curve, it's infrastructure. As someone who works in this space directly, the shift I've watched firsthand is from "AI that suggests a line of code" to "AI that owns an entire task end to end," tests, commits, and all. That changes what a developer's job actually is: less time typing syntax, more time on architecture, judgment calls, and reviewing what the agent produced. Businesses evaluating these tools should think less about which

one writes the prettiest code, and more about which one fits how their team actually wants to work: inside an editor, or from the terminal with full autonomy.

The Throughline

Look across all three industries and a pattern emerges: 2026 is the year AI stopped being evaluated on whether it *can* do the work, and started being evaluated on whether it can be *trusted* to do the work reliably, safely, and at scale. Enterprises are building governance alongside their agents. Regulators are building clearance pathways that keep pace with patient-facing AI. And developers are handing over more autonomy to coding agents precisely because the track record has started to earn it.

None of this means AI is finished maturing, far from it. But the businesses pulling ahead this month aren't the ones with the flashiest demo. They're the ones that figured out how to make AI boring, reliable, monitored, and simply part of how the work gets done.

THE APPLIED AI BRIEF · Issue #1 · August 2026

Sources: Gartner 2026 CIO and Technology Executive Survey · TechTimes · Oracle AI Blog · OpenAI · Bipartisan Policy Center · McGuireWoods · Shakudo Blog · Tech Insider · Nipralo Technologies · Developers Digest